

Kubernetes, often abbreviated as K8s, is a powerful open-source platform for automating the deployment, scaling, and management of containerized applications. Developed originally by Google, Kubernetes is now maintained by the Cloud Native Computing Foundation (CNCF). Below is a comprehensive guide covering the key concepts, architecture, and best practices for Kubernetes.

1. Introduction to Kubernetes

What is Kubernetes?

Kubernetes is a container orchestration platform that automates the deployment, scaling, and operations of application containers across clusters of machines. It abstracts the underlying infrastructure and provides a unified API for managing containerized applications.

**1.1 Key Features of Kubernetes

- **Automated Deployment and Scaling:** Automatically deploy, manage, and scale applications.
- **Self-Healing:** Automatically replaces failed containers and restarts applications.
- **Load Balancing:** Distributes traffic across containers.
- **Storage Orchestration:** Manages storage resources for applications.
- **Service Discovery:** Provides DNS names for services and load balances across them.
- **Configuration Management:** Manages configuration and secrets for applications.
- **Rolling Updates:** Performs rolling updates to applications with zero downtime.

**1.2 Basic Terminology

Term	**Description**	
-----	-----	
Cluster	A set of nodes (physical or virtual machines) running Kubernetes components.	
Node	A machine (VM or physical) that runs Kubernetes components like the Kubelet	
and Docker.		

Pod	The smallest deployable unit that can contain one or more containers.	
Deployment	A higher-level abstraction for managing and scaling a set of Pods.	
Service	Defines how to access Pods, providing load balancing and service discovery.	
Namespace	Virtual clusters within a Kubernetes cluster to organize resources.	
ReplicaSet	Ensures a specified number of pod replicas are running at any given time.	
StatefulSet	Manages stateful applications, maintaining unique identities for Pods.	
DaemonSet	Ensures a copy of a Pod runs on all (or some) Nodes.	
Job	Manages the execution of one-time tasks and ensures completion.	
CronJob	Creates Jobs on a scheduled time basis.	
ConfigMap	Manages configuration data for applications.	
Secret	Stores sensitive data, such as passwords or tokens.	
Ingress	Manages external access to services, typically via HTTP/HTTPS.	
PersistentVolume (PV)	A piece of storage in the cluster provisioned by an administrator.	
PersistentVolumeClaim (PVC)	A request for storage by a user.	
Helm	A package manager for Kubernetes applications.	

2. Kubernetes Architecture

2.1 Kubernetes Components

Component	Description	
-----	-----	

Master Node	Manages the Kubernetes cluster. Contains components like API server, controller manager, and scheduler.
Kubelet	An agent that runs on each worker node, ensuring that containers are running in Pods.
Kube-Proxy	Maintains network rules for Pod communication and load balancing.
API Server	The entry point for all REST commands used to control the cluster.
Controller Manager	Ensures that the desired state of the cluster is maintained.
Scheduler	Assigns Pods to Nodes based on resource availability and constraints.
etcd	A distributed key-value store used for storing all Kubernetes cluster data.

2.2 Kubernetes Object Lifecycle

Lifecycle Stage	Description	
-----	-----	
Creation	Define objects using YAML or JSON manifests and apply them using `kubectl apply`.	
Update	Modify the configuration and apply changes using `kubectl apply`.	
Scaling	Adjust the number of Pods using `kubectl scale`.	
Deletion	Remove objects using `kubectl delete`.	

Diagram of Kubernetes Architecture:

![[Kubernetes Architecture]](<https://www.redhat.com/cms/managed-files/k8s-architecture-2020-09-17.jpg>)

Source: Red Hat

3. Core Kubernetes Concepts

3.1 Pods

- **What is a Pod?**

- A Pod is the smallest and simplest Kubernetes object. A Pod encapsulates one or more containers.

- **Example YAML for a Pod:**

```
```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-pod
spec:
 containers:
 - name: my-container
 image: nginx:latest
 ports:
 - containerPort: 80
```
```

3.2 Deployments

- **What is a Deployment?**

- A higher-level abstraction for managing a set of Pods. It ensures the desired state is maintained.

- **Example YAML for a Deployment:**

```
```yaml
apiVersion: apps/v1
kind: Deployment
```

```
metadata:
 name: my-deployment
spec:
 replicas: 3
 selector:
 matchLabels:
 app: my-app
 template:
 metadata:
 labels:
 app: my-app
 spec:
 containers:
 - name: my-container
 image: nginx:latest
 ports:
 - containerPort: 80
 ...
```

### **\*\*3.3 Services\*\***

- **\*\*What is a Service?\*\***
- Provides a stable IP address and DNS name for a set of Pods.
- **\*\*Example YAML for a Service:\*\***

```
``yaml
apiVersion: v1
kind: Service
metadata:
 name: my-service
spec:
```

```
selector:

 app: my-app

ports:
 - protocol: TCP
 port: 80
 targetPort: 80
...

```

### **\*\*3.4 ConfigMaps and Secrets\*\***

- **\*\*ConfigMaps\*\***: Manage configuration data for applications.
- **\*\*Secrets\*\***: Store sensitive data such as passwords and tokens.
- **\*\*Example YAML for a ConfigMap:\*\***

```
``yaml
apiVersion: v1
kind: ConfigMap
metadata:
 name: my-config
data:
 key1: value1
 key2: value2
...

```

- **\*\*Example YAML for a Secret:\*\***

```
``yaml
apiVersion: v1
kind: Secret
metadata:
 name: my-secret

```

```
type: Opaque
data:
 username: dXNlcg== # Base64 encoded username
 password: cGFzc3dvcmQ= # Base64 encoded password
...
```

### **\*\*3.5 Persistent Storage\*\***

- **\*\*PersistentVolume (PV)\*\***: A storage resource in the cluster.
- **\*\*PersistentVolumeClaim (PVC)\*\***: A request for storage.
- **\*\*Example YAML for a PV:\*\***

```
```yaml
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
  hostPath:
    path: "/mnt/data"
...```
```

- ****Example YAML for a PVC:****

```
```yaml
apiVersion: v1
```

kind: PersistentVolumeClaim

metadata:

name: my-pvc

spec:

accessModes:

- ReadWriteOnce

resources:

requests:

storage: 1Gi

...

### **\*\*3.6 Ingress\*\***

- **\*\*What is Ingress?\*\***

- Manages external access to services, typically HTTP/HTTPS.

- **\*\*Example YAML for an Ingress:\*\***

```
```yaml
```

apiVersion: networking.k8s.io/v1

kind: Ingress

metadata:

name: my-ingress

spec:

rules:

- host: myapp.example.com

http:

paths:

- path: /

pathType: Prefix

backend:

service:


```
    name: my-service

    port:

    number: 80
...

```

****3.7 Helm Charts****

- **What is Helm?**

- A package manager for Kubernetes, similar to apt for Debian-based systems or yum for Red Hat-based systems.

- **Basic Commands:**

- **Install a Chart:**

```
```bash
helm install my-release stable/nginx
...

```

##### **- \*\*Upgrade a Release:\*\***

```
```bash
helm upgrade my-release stable/nginx
...

```

- **Uninstall a Release:**

```
```bash
helm uninstall my-release
...

```

### **\*\*Helm Chart Example:\*\***

```
```yaml
apiVersion: v2
name: mychart
description: A Helm chart for Kubernetes
version: 0.1.0

```

dependencies:

- name: nginx

version: 1.16.0

repository: <https://charts.bitnami.com/bitnami>

templates:

- name: deployment.yaml

apiVersion: apps/v1

kind: Deployment

spec:

replicas: 3

selector:

matchLabels:

app: my-app

template:

metadata:

labels:

app: my-app

spec:

containers: